

# Building An Accurate And Efficient Animal Classification Application

A Deep Learning Project

Jingyi Qiu  
Toyesha Kaushik

# Index

Introduction	2
Model Enhancement	2
Conclusion	3
Appendices	4

# Introduction

Our goal is to develop an efficient and accurate animal image identification algorithm. This algorithm will be integrated into an application that is accurate, compact, and fast to run on any device, including PCs, smartphones, and smart watches. Our focus will be on optimizing both accuracy and efficiency to identify the best-performing algorithm. We will evaluate and modify a baseline model as needed to meet our specifications and achieve optimal performance.

Accuracy is determined by using a function that computes the accuracy over the top-k predictions for the specified values of k. For example, in top-3 accuracy the output is considered the ‘right answer’ if the right answer appears in the top 3 guesses. Meanwhile, efficiency is determined using FLOPS (Floating-Point Operations per Second). This is a unit of measurement used to quantify the computing power (Sheldon, R. 2023). Efficiency is calculated by dividing accuracy by FLOPS count.

# Model Enhancement

## Data Preprocessing

Before the algorithm can be created, the quality of collected data, the number of samples per class (data distribution), and the image type should be thoroughly examined. Some images, despite being high quality, might have caused problems e.g clear reflections look like double-headed animals. Some images, such as drawings, were left in the dataset if accurate. Some subjects were far too camouflaged to be reliably used as training data. Most images, due to legal reasons, had watermarks. While most of them could be ignored, cases where the watermarks were too obstructive needed to be removed.

See *Appendix A* for some more examples.

## Baseline Model Evaluation

Base models 1 and 2 use the original provided code, each containing 857,239 training parameters. Note that model 1 utilizes the raw, unprocessed dataset whereas model 2 is trained on the preprocessed data. However, as shown in the Ablation Study below, the accuracy barely increases, with efficiency remaining below 60%. The Loss-Epochs graphs (see

*Appendix B*) show a significant overfitting in the models. This could be due to the CNN model architecture, chosen learning rate, optimizer or number of epochs.

## Ablation Study

Cross-entropy is consistently applied to all models. It is the most suitable lost function for multiclass classification problems (Andreieva, V. 2021), whereas other loss functions are more reliable for binary classification problems.

Model no.	Max Epochs	BN + Dropout	Optimizer	NN details	LR	Accuracy (%)	FLOPS (G)	Efficiency	Comments
1. (baseline)	10	N	Adam	CNN	0.001	37.86	0.69	54.87	Overfitting
2. (baseline)	10	N	Adam	CNN	0.001	40.25	0.69	58.33	Overfitting
3.	15	N	NAdam	CNN	0.001	37.17	0.69	53.87	Overfitting, worse acc.
4.	50	N	SGD	CNN	0.01	39.76	0.69	57.62	Overfitting
5.	15	Y - 0.5	Adam	CNN	0.001	58.31	0.69	84.51	Less overfit Increased acc.
6.	50	Y - 0.5	Adam	CNN + DA (RandomRotation, RandomResizedCrop)	0.001	65.05	0.69	94.28	Less overfit
7.	15	Y - 0.5	Adam	CNN, Greyscale	0.001	37.62	0.62	60.68	Less overfit
8.	15	N	Adam	Transfer Learning, resnet18, RGB	0.001	75.32	0.96	78.46	Trade-Off considered
9.	15	N	Adam	Transfer Learning, diff. pre training data	0.001	-	-	-	Insufficient processing power

## Evaluation and Visualization

**Epochs:** Started with 10 and increased to 50. This increase clearly exacerbated the overfitting tendencies of the SGD optimizer in Model 4. Therefore, the value was adjusted to 15 to achieve a balance between computational time and sufficient data.

**Optimizers:** The best optimizer returning highest accuracy, as gathered from models 2, 3, and 4, is Adam, and was therefore used to train the remaining models.

**Learning rate (LR):** Increasing from 0.001 to 0.01 led to significant gradient overshooting during training in Model 4. Hence, LR=0.001 was the most optimal value.

**CNN architecture:** Batch Normalization and Dropout layers were introduced in order to reduce the effects of overfitting. Results are immediate. Accuracy and efficiency increased significantly from Model 4 to Model 5 after implementing BN + dropout.

**Data augmentation:** Advanced techniques, such as RandomRotation and RandomResizedCrop, were introduced to expand the variety of training data and minimize overfitting. Accuracy & Efficiency increased further from Model 5 to Model 6.

**Color channels:** Retraining Model 6 on grayscale images decreased the accuracy significantly, possibly due to color being a key feature in differentiating animals of a similar species, e.g. birds.

Transfer learning: Most balanced performance in terms of accuracy/efficiency tradeoff.

See *Appendix B* for detailed visualization for all models.

## Conclusion

### Model Selection

Realistically, an animal classification program available to the general public would be more popular with a high accuracy. Therefore, the most optimal model is Model 8, which uses Transfer Learning to achieve this. The efficiency is also at an acceptable level, close to 80%. A tradeoff between accuracy and efficiency was favored, as efficiency may vary depending on the device being used.

### Limitations

As seen in *Appendix D*, Google Colab GPU and RAM overloaded, and therefore local ones were used after Models 4 and 6 respectively. This may lead to minor discrepancies in expected accuracy values.

The main challenge faced when training all models, including our most efficient one, is creator bias in choosing the training data images. We found that:

1. Some images in the same class may be too similar, leading to overfitting
2. Choosing images to use is a highly objective process. What may be considered as a “good” image may not be considered one by another contributor
3. Limited knowledge could lead to data contamination. For example, *Panthera onca* and *Panthera pardus* look extremely similar to the naked eye, so an image of one could accidentally be classified as another.
4. Unequal numbers of training images for each class could lead to identification bias.

There is also difficulty in sourcing images. This could be because:

1. There is a lack of representation of sex of an animal, leading to difficulty in identification. E.g. *Panthera leo* males are considered the ‘default’, leading to less images of females, which may mean more difficult identification
2. Lack of representation of life stages. E.g. Pupal/larval stages are not frequently present in datasets
3. Most images available on a public domain are watermarked. While some watermarks are less obvious and can be ignored, some cover the subject of the image, and therefore cannot be used.
4. Some animals are difficult to photograph. This could either be due to ease of access e.g. underwater animals such as *Monodon monoceros*, or due to extinction of certain classes.
5. Drawings or models of animals may not be scientifically accurate, leading to discrepancies.

Another possible issue is training the model to ignore certain aspects of a background or include the background. For example, the *Pavo cristatus* has feathers that may be ignored due to being considered a

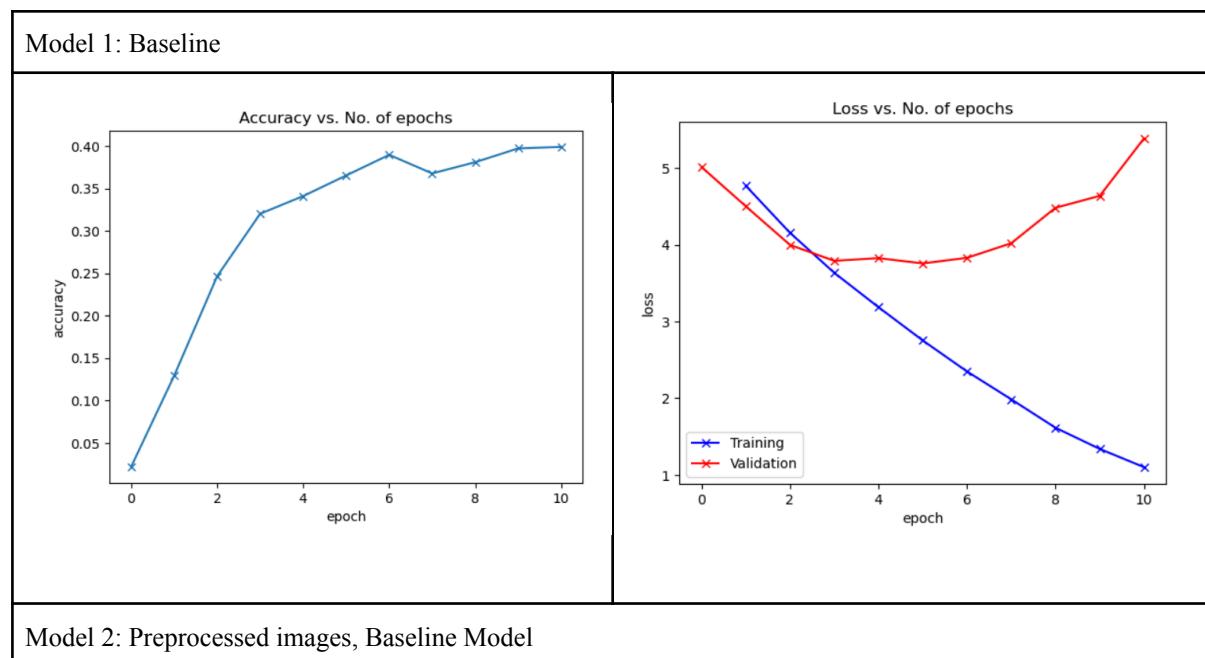
background, whereas the *Lemur catta* images, due to the frequent presence of branches, may have the branches mistakenly identified as key features.

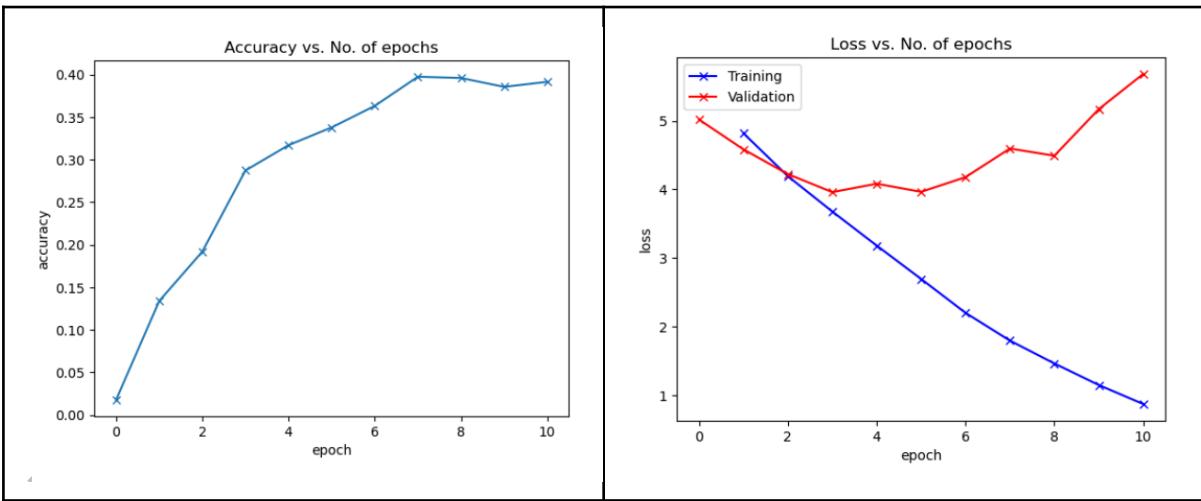
## Appendix A - Dataset preprocessing

Here are some examples of the images that needed to be disregarded in the dataset:

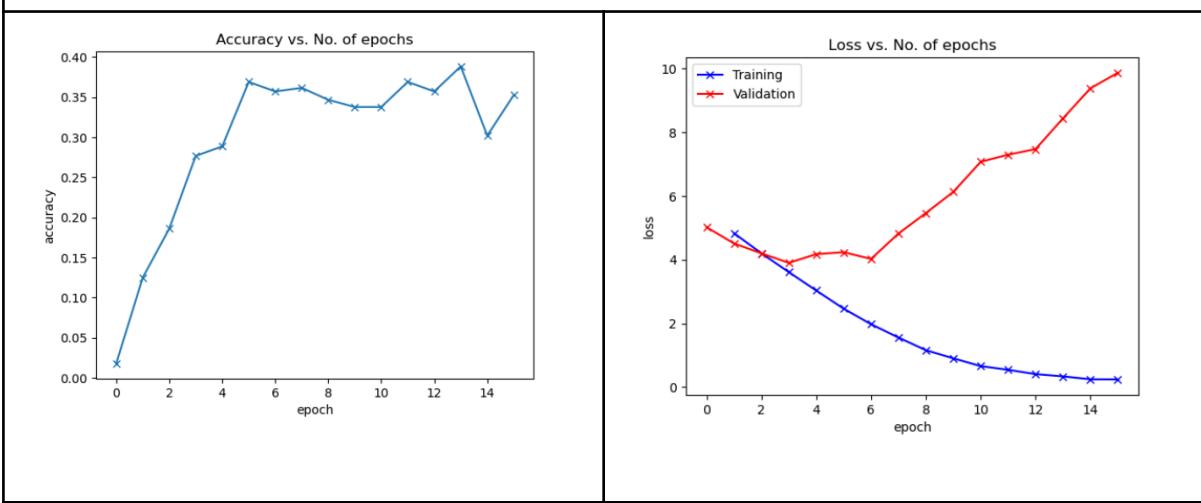
Class	Original Sample Size	Problems	Final Sample Size	Quality
<i>aethia-cristatella</i>	50	Fake/Seemingly AI-generated Too compressed Subject too far away Odd angle Overlapping subjects Labels covering subject	46	High
<i>agalychnis-callidryas</i>	34	Other objects as the focal point Zoomed in too far to see distinguishable features Odd angle Uncommon posture	29	High
<i>ailuropoda-melanoleuca</i>	31	Other animals in the picture Albinism (rare case/misclassification)	28	Mid

## Appendix B - Visualization

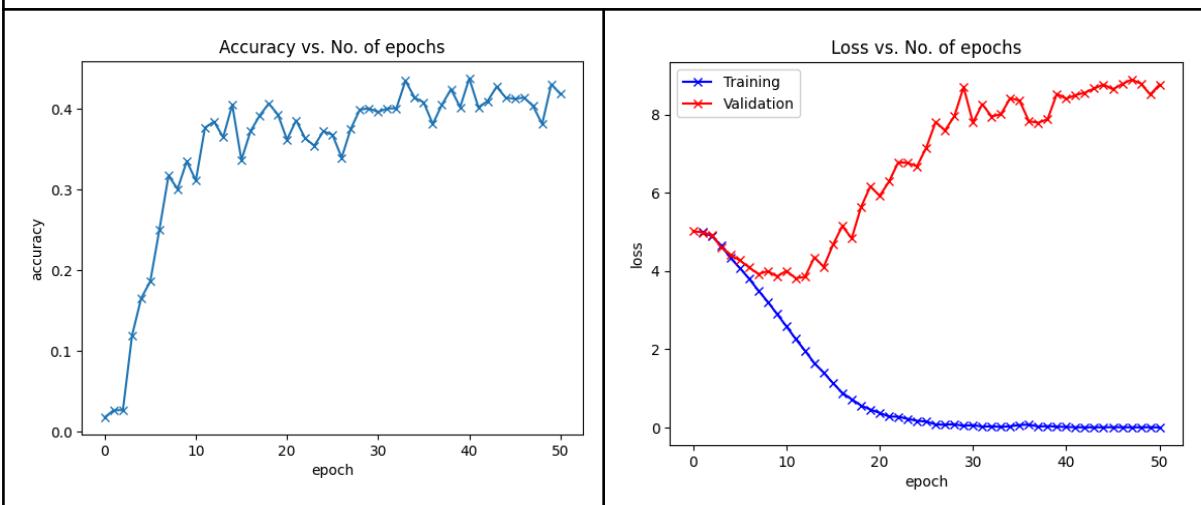




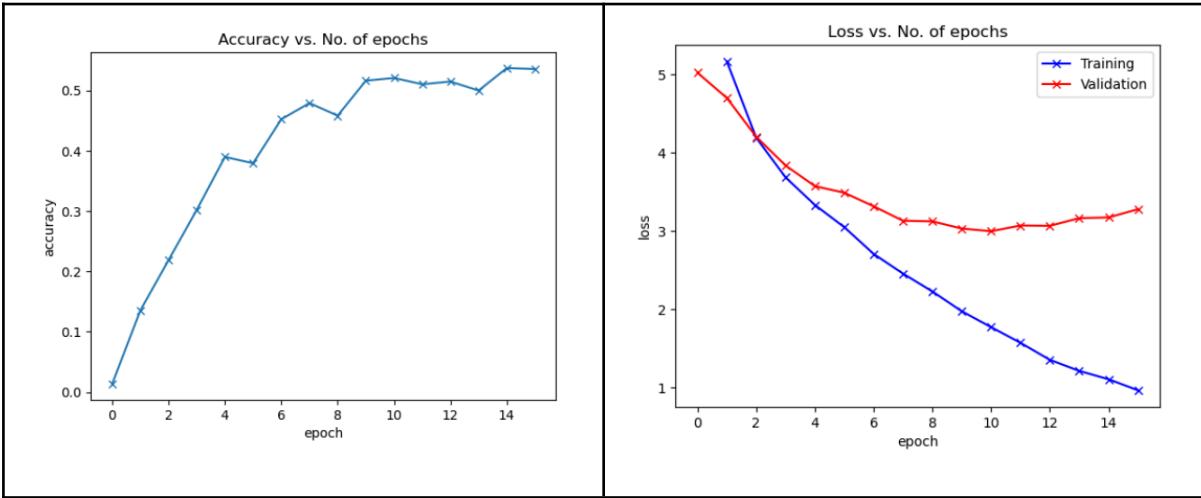
Model 3: NAdam



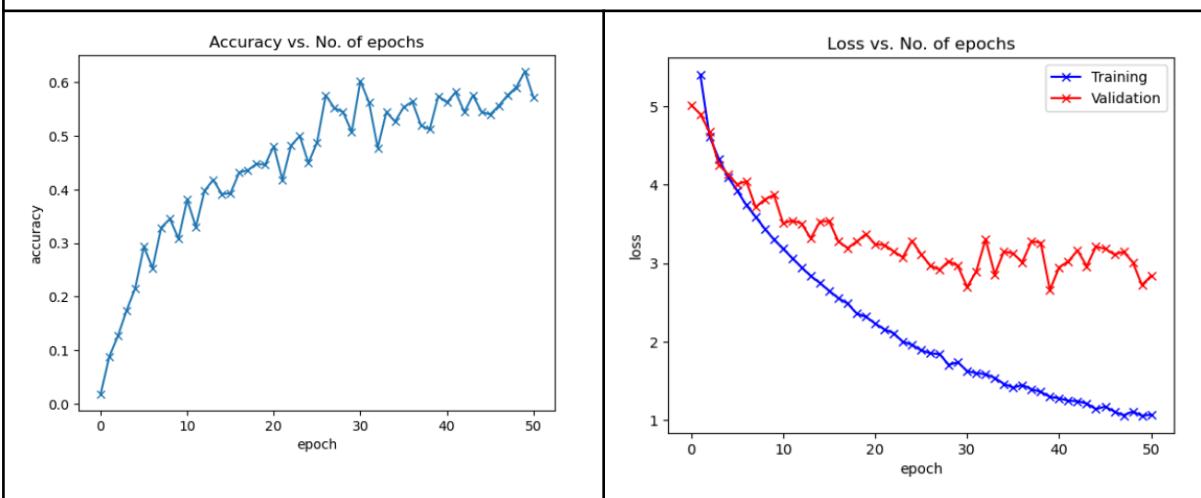
Model 4: SGD



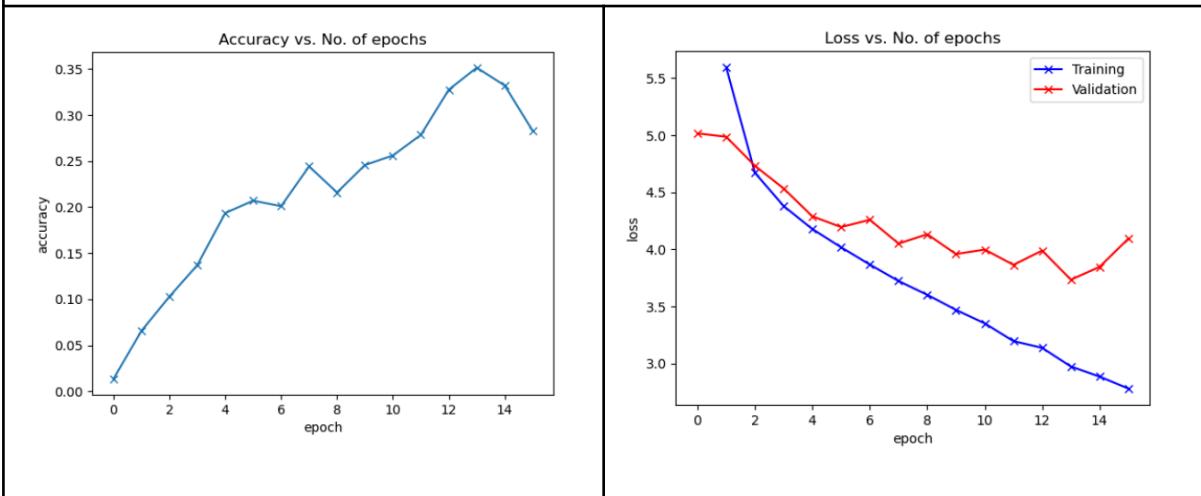
Model 5: BN, Dropout



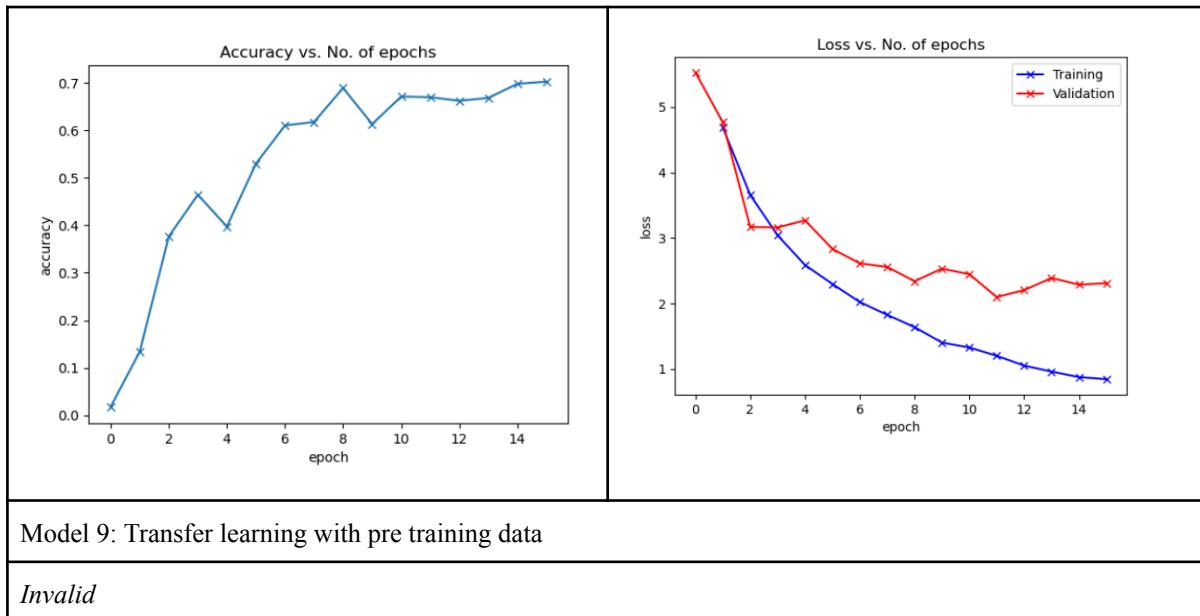
Model 6: Data Augmentation



Model 7: Greyscale



Model 8: Transfer Learning (**FINAL/MOST OPTIMAL**)



## Appendix C - References

Andreieva, V. et al, *Generalization of Cross-Entropy Loss Function for Image Classification* (2021, Jan). Published via Mohyla Mathematical Journal, accessed via ResearchGate [19 Jun 2024] <<https://www.researchgate.net/publication/349850933> Generalization of Cross-Entropy Loss Function for Image Classification>

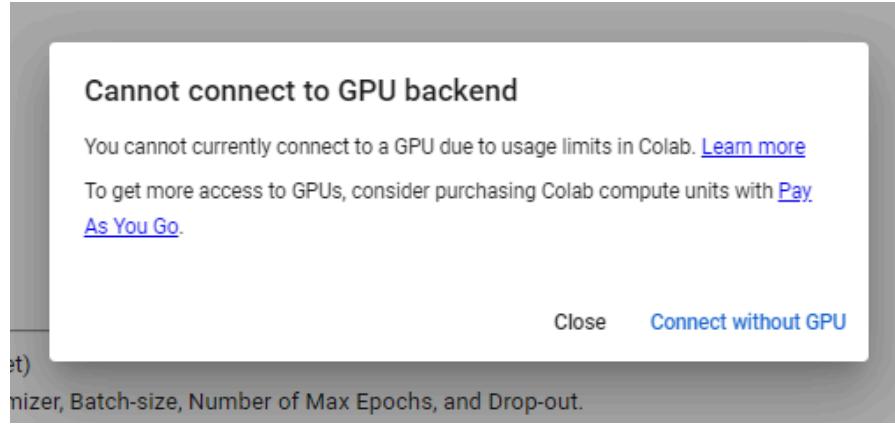
Sheldon, R. *floating-point operations per second (FLOPS)* (2023, Aug). Published and accessed via TechTarget [19 Jun 2024] <<https://www.techtarget.com/whatis/definition/FLOPS-floating-point-operations-per-second>>

## Appendix D - GPU & RAM limits

The screenshot shows a Jupyter Notebook interface with the following details:

- Title:** Animal\_Classification.ipynb
- Content:**
  - Table of contents:** Animal Competition (15%)
  - Code cell:**
    - Code block: 

```
Indented block
```
    - Text: "For the non-competition mode, we will use the Animal (<https://cloudstor.aarnet.edu.au/plus/s/c2YtNAevhWD6uBx>) dataset. This dataset contains images of 151 different animals."
    - Text: "The dataset contains a total of 6270 images corresponding to the name of animal types. All images are RGB images of 112 pixels wide by 112 pixels high in jpg format. The images are separated in 151 folders according to their respective class."
    - Text: "The task is to categorize each animal into one of 151 categories. We provide baseline code that includes the following features:
      - Loading and Analysing the dataset using torchvision.
      - Defining a simple convolutional neural network.
      - How to use existing loss function for the model learning.
      - Train the network on the training data.
      - Test the trained network on the testing data.    - Text: "The following changes could be considered:
      - 1. "Transfer" Learning (ie use a model pre-trained another dataset)
      - 2. Change of advanced training parameters: Learning Rate, Optimizer, Batch-size, Number of Max Epochs, and Drop-out.
      - 3. Use of a new loss function.
      - 4. Data augmentation
      - 5. Architectural Changes: Batch Normalization, Residual layers, etc.
      - 6. Others    - Text: "Your code should be modified from the provided baseline. A pdf report of a maximum of two pages is required to explain the changes you made from the baseline, why you chose those changes, and the improvements they achieved."
    - Marking Rules:** "We will mark the competition based on the final test accuracy on testing images and your report. Final mark (out of 50) = acc\_mark + efficiency mark + report mark"
    - Acc\_mark 10:** "We will rank all the submission results based on their test accuracy. Zero improvement over the baseline yields 0 marks. Maximum 10 marks."
  - Message:** "Cannot connect to GPU backend" (with a link to "Learn more" and "Pay As You Go").
  - Runtime Status:** "Allocating runtime"



The screenshot shows the same "Animal\_Classification.ipynb" notebook in Google Colab. The code cell contains PyTorch code for transfer learning. The output of the cell shows training progress and a final evaluation step. A progress bar indicates the final evaluation took 1.260ms. The session then crashes, as indicated by the message "Your session crashed after using all available RAM." at the bottom of the screen. The status bar at the bottom right shows "Connected to Python 3 Google Compute Engine backend".

```
from torchvision import models

class TransferLearningModel(ImageClassificationBase):
    def __init__(self, num_classes):
        super().__init__()
        self.model = models.resnet18(pretrained=True)
        in_features = self.model.fc.in_features
        self.model.fc = nn.Linear(in_features, num_classes)

    def forward(self, x):
        return self.model(x)

# Create the model
transfer_model = TransferLearningModel(num_classes)
transfer_model = to_device(transfer_model, device)

# Train the model with transfer learning
transfer_history = [evaluate(transfer_model, val_loader)]
transfer_history

[{'val_loss': 5.6967597007751465, 'val_acc': 0.0223214291036129}]

[ ] transfer_history += fit(num_epochs, lr, transfer_model, train_dl, val_dl, opt_func)

100% 235/235 [08:32-00:00. 1.91ms]
Epoch [0], train_loss: 4.15469, val_loss: 4.1574, val_acc: 0.2262
80% 189/235 [02:36-00:36. 1.260ms]

[ ] plot_accuracies(transfer_history)

[ ] plot_losses(transfer_history)

[ ] # Evaluate the transfer learning model
evaluate(transfer_model, test_loader)

Your session crashed after using all available RAM. X LOE
```